

7.3 Skip-gram Word Embeddings

Edwin Simpson

Department of Computer Science,
University of Bristol, UK.

Dense Vectors

- Term-document matrices are sparse with large numbers of dimensions and many near-zero values;
- A dense vector representation of 50-1000 values has advantages:
 - Fewer dimensions means models for tasks such as classification and sequence labelling need fewer parameters;
 - This leads to less **over-fitting** and better generalisation;
 - Relations such as synonymy can be better represented.
- Dense vectors can be learned using the **skip-gram** model,
 - Implemented by the software **word2vec**;
 - Alternatives: continuous bag of words; GloVe.

Skip-gram: Core Ideas

- Use the **contextual** view of meaning
 - Distributional hypothesis
 - Determine a word's meaning from its neighbouring words
- Represent the context that a target word t occurs in:
 - Term-document matrix: whole document
 - Skip-gram: **context** window of $\pm k$ words either side of the target word

```
... lemon,    a [tablespoon of apricot jam,    a] pinch ...
           c1          c2      t      c3          c4
```

Skip-gram: Core Ideas

- Embeddings as a by-product of a classifier:
 - Learn a classifier to distinguish real and fake contexts for any given t
 - The parameters of this classifier form an **embedding vector**
- **Self-supervised learning:**
 - No training labels, learn to predict part of the text itself
 - Positive examples: the real context words in a set of training documents
 - Negative examples: randomly sampled words from the vocabulary.

Prediction with Skip-gram


$$P(+ \mid \mathbf{c}, \mathbf{t}) = \prod_{i=1}^k \frac{1}{1 + e^{-\mathbf{t} \cdot \mathbf{c}_i}}$$

Prediction with Skip-gram

- Recall that cosine similarity is a normalised dot product;
- We combine the embedding vectors for target word t and context word c_i using the dot product, $t \cdot c_i$.
- The value of $t \cdot c_i$ ranges from $-\infty$ to ∞ .

Prediction with Skip-gram

- **Sigmoid function** maps real values to numbers between 0 and 1.
- $\sigma(\mathbf{t} \cdot \mathbf{c}_i) = \frac{1}{1+e^{-\mathbf{t} \cdot \mathbf{c}_i}}$
- This means we have a **logistic regression** classifier for each target word with weights \mathbf{t}
- It maps input vectors \mathbf{c}_i to probabilities.



Prediction with Skip-gram

- If a context \mathbf{c} is true (positive), the occurrences of all its individual words must be positive;
- Assume the truth values of context words are **conditionally independent**:

$$P(+ \mid \mathbf{c}, \mathbf{t}) = \prod_{i=1}^k \frac{1}{1 + e^{-\mathbf{t} \cdot \mathbf{c}_i}}$$

- Summary:
 - Combine embeddings of target and context words using dot product;
 - Apply sigmoid function to obtain a probability;
 - Take a product of independent probabilities.

Choosing Negative Examples

- Noise words are selected randomly according to their weighted frequency:

$$P_{\alpha(w)} = \frac{\text{count}(w)^\alpha}{\sum_{w'} \text{count}(w')^\alpha}$$


- Selecting by frequency avoids comparing with rare negative words that are very easy to spot as fake;
- Weighting the frequencies by 0.75 is a heuristic that prevents selecting too many very frequent words like 'the'.

Learning Objective

- Goal: choose \mathbf{t} and \mathbf{c}_i to minimise prediction error on the training set
- $L(\mathbf{T}, \mathbf{C}) = -\sum_{(t,c) \in +} \log P(+|\mathbf{t}, \mathbf{c}) - \sum_{(t,c) \in -} \log P(-|\mathbf{t}, \mathbf{c})$
- Two sets of embeddings for each word!
 - Target word embeddings;
 - Context word embeddings.
 - Typically, only the target embeddings \mathbf{T} are taken from the model for use as embeddings in downstream tasks.

Gradient Descent

- Consider each dimension of each target and context vector as a 'weight', w .
- Starting from random initial values of w ...
- Increase or decrease w in the opposite direction to the gradient of $L(\mathbf{T}, \mathbf{C})$:
- Thereby reduce the loss.



Mini-batch Stochastic Gradient Descent

$$w^{t+1} = w^t - \frac{\eta}{N} \nabla_w L(\mathbf{T}, \mathbf{C})$$

- $L(\mathbf{T}, \mathbf{C})$ is a sum over contexts, so is $\eta \nabla_w L(\mathbf{T}, \mathbf{C})$.
- So, split the training data into **batches**, compute the gradient for one batch at a time:

$$w^{t+1} = w^t - \frac{\eta}{B} \nabla_w L(\mathbf{T}_i, \mathbf{C}_i)$$

- Faster updates, parallelisation

∇_w : the gradient with respect to weight w .

N : number of training contexts.

η : learning rate that controls the size of each step.

B : batch size.

Mini-batch Stochastic Gradient Descent

- To learn the skipgram embeddings:
 - Iterate over batches;
 - Within each batch, iterate over the ‘weights’, i.e., the parameters in the target word and context vectors;
 - Output target word vectors as embeddings.
- Same algorithm is used for:
 - **Logistic regression**
 - **Neural networks**

Summary

- Skip-gram applies the distributional hypothesis to learn dense vector representations (embeddings) of words from their contexts;
- We learn a binary classifier that distinguishes real (+ve) and fake (-ve) contexts for any given target word;
- Stochastic gradient descent trains the classifier by iteratively making adjustments to the weights to reduce the loss
- The weights of the classifier for a particular target word are used as its embedding.
- Words with similar contexts will have similar embeddings.