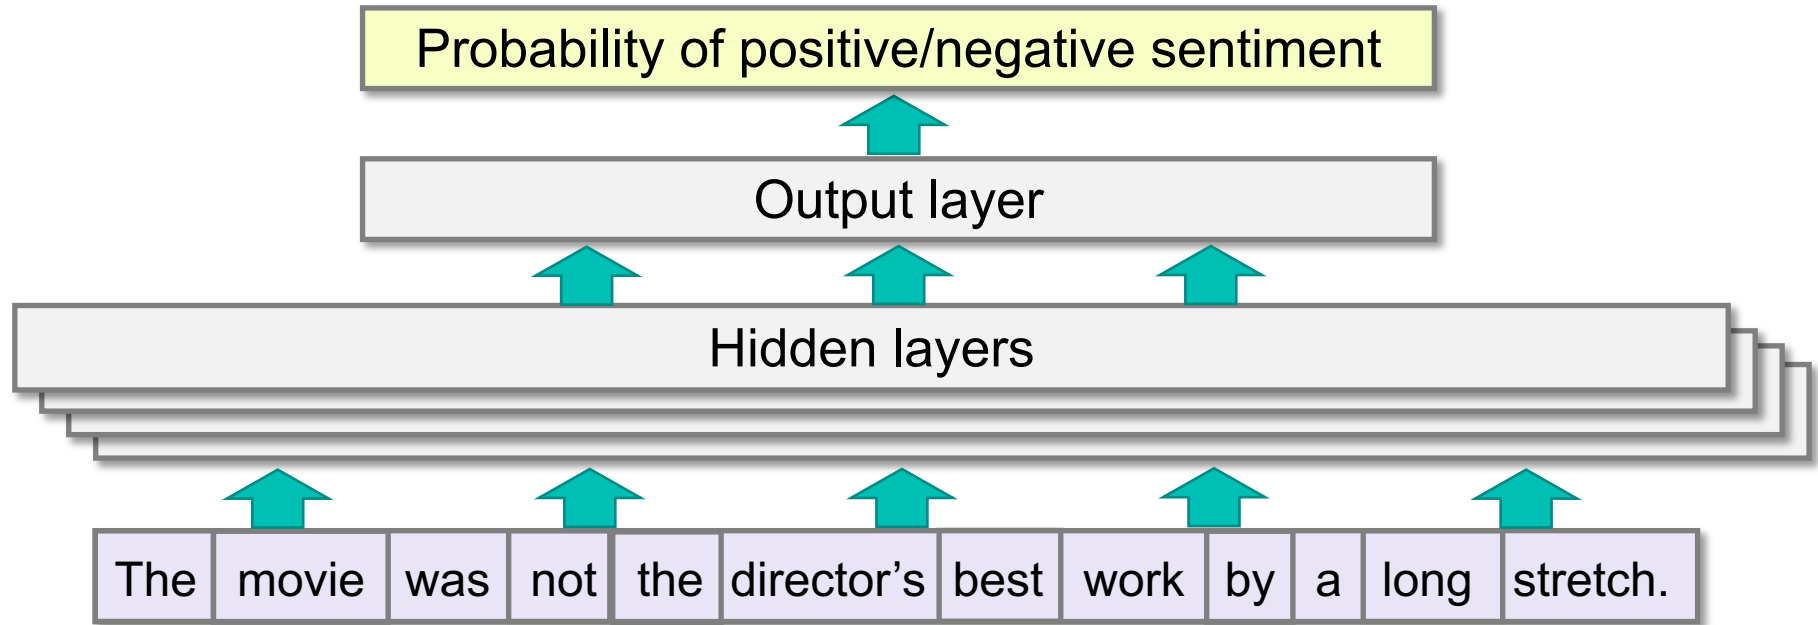


# 8.1 Deep Text Classifiers

Edwin Simpson

Department of Computer Science,  
University of Bristol, UK.

# Neural Network (NN) Text Classifier

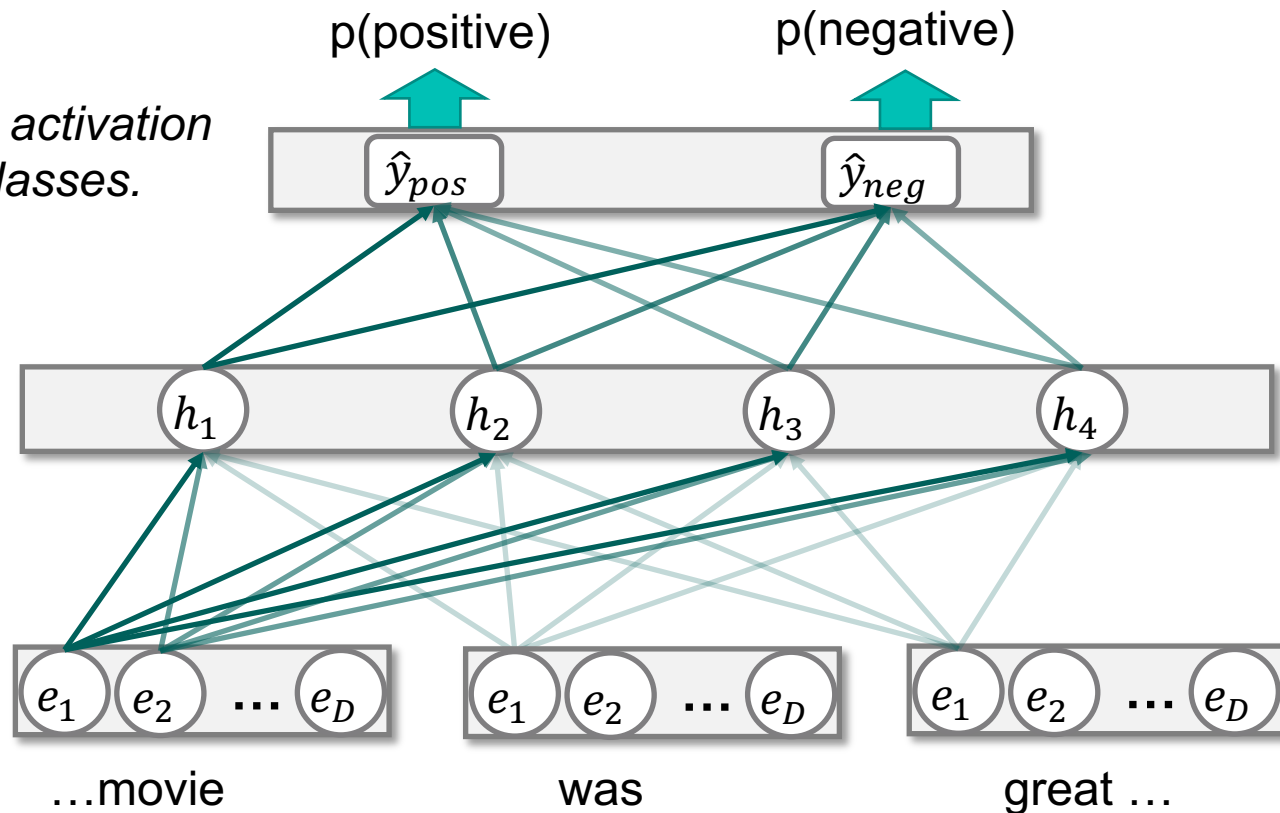


# Neural Text Classifier

*Output layer: softmax activation function for multiple classes.*

*Fully-connected Hidden layer: ReLU activation function.*

*Input layer: embeddings or one-hot encoding*



# Fine-tuning Embeddings

- Word embeddings can be stored in a  $|V| \times D$  matrix,  $E$ :
- Look up a word's embedding by multiplying its one-hot encoding with the matrix.
- $E$ , can be treated as another weight vector in an embedding layer.
- Then we can train the embedding matrix through backpropagation.
  - Embeddings are **pretrained** using a model like Skipgram;
  - Then **fine-tuned** on the training set to better represent the features that are important for this task.
  - This is **end-to-end** deep learning: the whole pipeline sits inside one model.

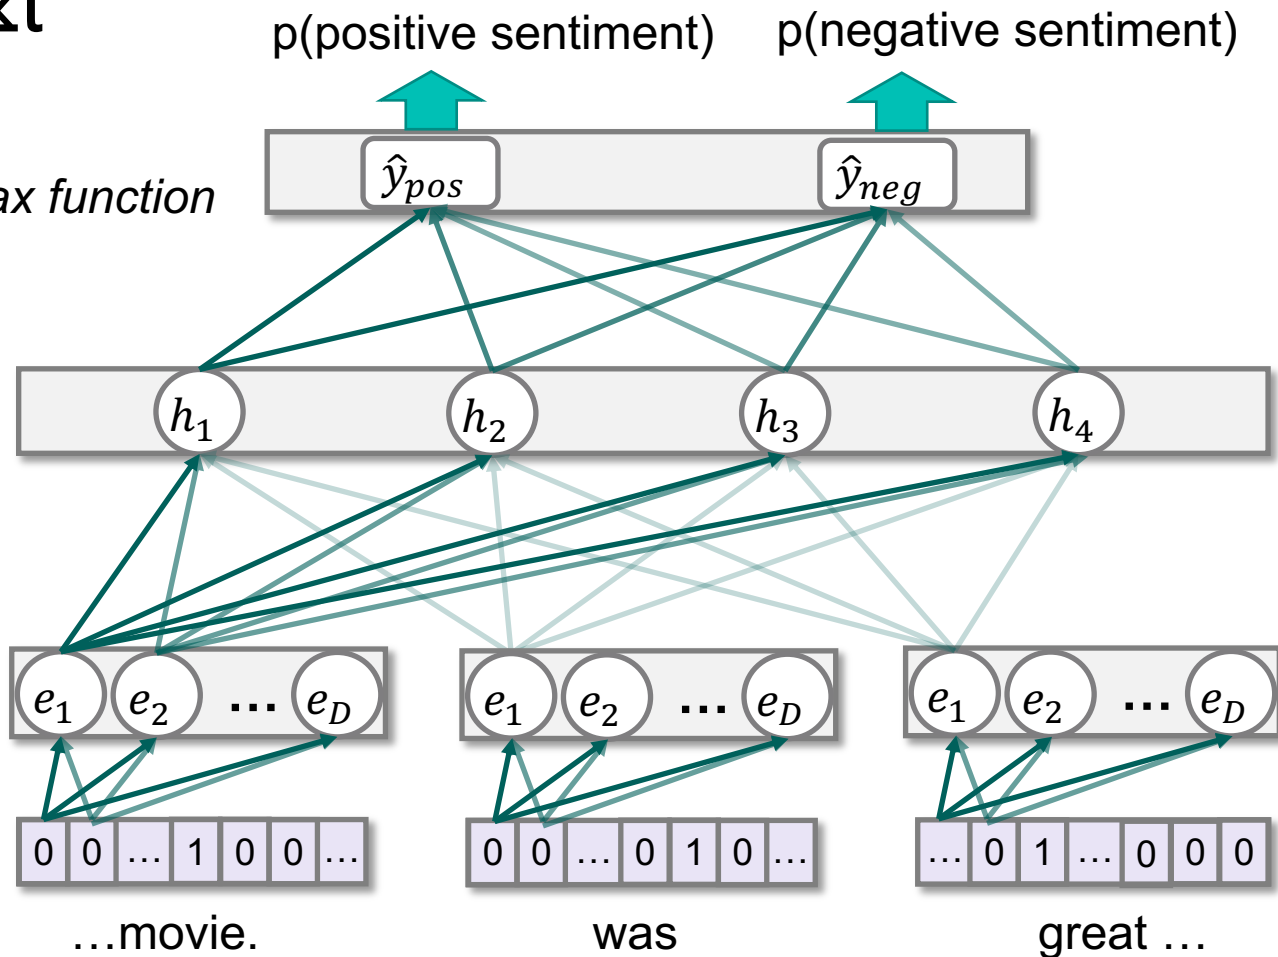
# Neural Text Classifier

*Output layer: softmax function*

*Feed-forward  
Hidden layer*

*Embedding layer*

*Input layer:  
one-hot encoding*



# Padding

- The weight matrix  $W^{(l)}$  for each layer has a fixed size.
- But the input texts can be of varying lengths.
- Define a maximum document length  $N$  and truncate long documents.
- **Pad** short documents by appending special '[PAD]' tokens to the sequence until it has length  $N$ .

# Summary

- Neural networks can process sequences of words for tasks like text classification.
- Inputs can be either one-hot encodings or fixed word embeddings.
- If we include an embedding layer, the embeddings can be fine-tuned to the training set.
- Padding is needed to convert documents of varying lengths to a fixed-size input.