

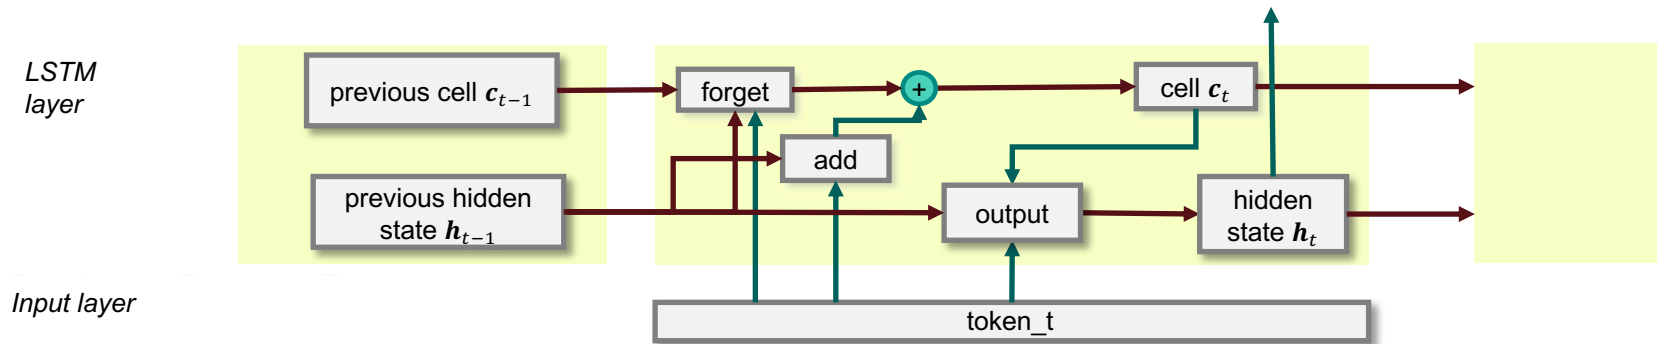
9.1: Self-Attention & Transformers

Edwin Simpson

Department of Computer Science,
University of Bristol, UK.

Recap: LSTMs

- Sequential text processing allows a model to interpret each token within the context of a sentence and compose the meaning of the sequence from the individual words.
- LSTMs are a type of RNN for sequential text processing.
- The LSTM memory cell stores context information that passes along the sequence to help make the decision for each token.



Drawbacks of LSTMs

- **Scalability:** processing each token in turn means we cannot parallelise the processing of a sequence.
- **Long-range dependencies:**
 - LSTMs retain information over a longer distance than standard RNNs;
 - But context information still has to pass through each time-step until it's needed.
 - Small changes to the cell mean information is still lost and gradients go to zero over long sequences.
- **Limited capacity of the memory cell:**
 - Complex sentences may require the memory to retain multiple pieces of information;
 - But a single memory cell may forget earlier context that becomes important later on.

Connections in Neural Networks

- The drawbacks of LSTMs stem from the recurrent connections that pass information along the sequence.
- We can improve this with a new type of connection!
- So far we have encountered two types of neural network layers that use different kinds of connection:

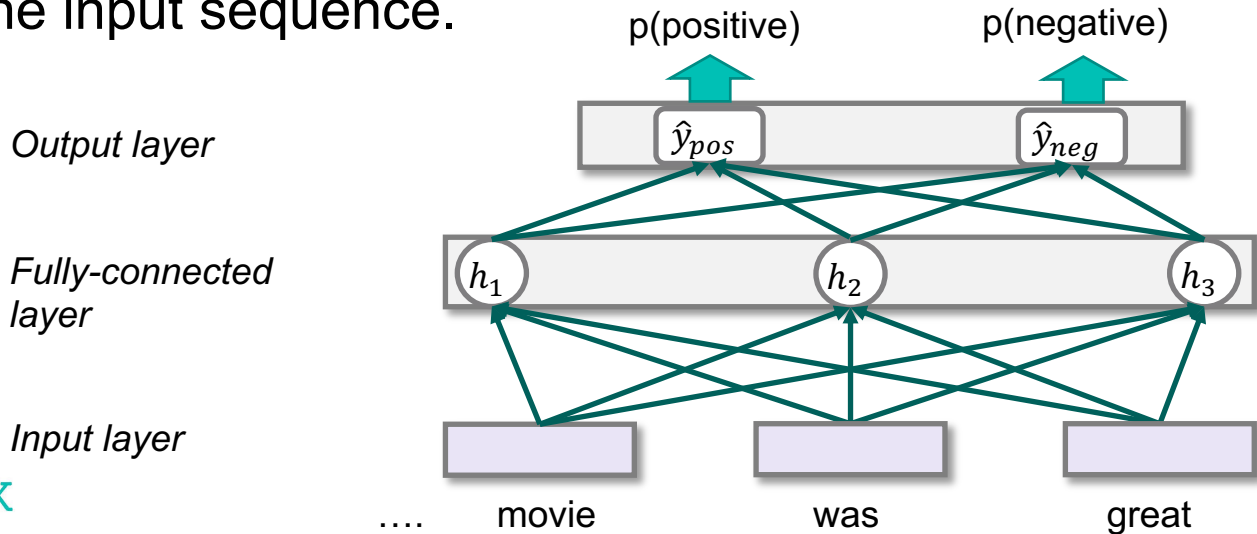
Which two types of layers have we looked at so far?

Connections in Neural Networks

- The drawbacks of LSTMs stem from the recurrent connections that pass information along the sequence.
- We can improve this with a new type of connection!
- So far we have encountered two types of neural network layers that use different kinds of connection:
 - Fully-connected layers with feed-forward connections.
 - Recurrent layers with recurrent connections.
- The connections in both types of layer are fixed when we design the model.

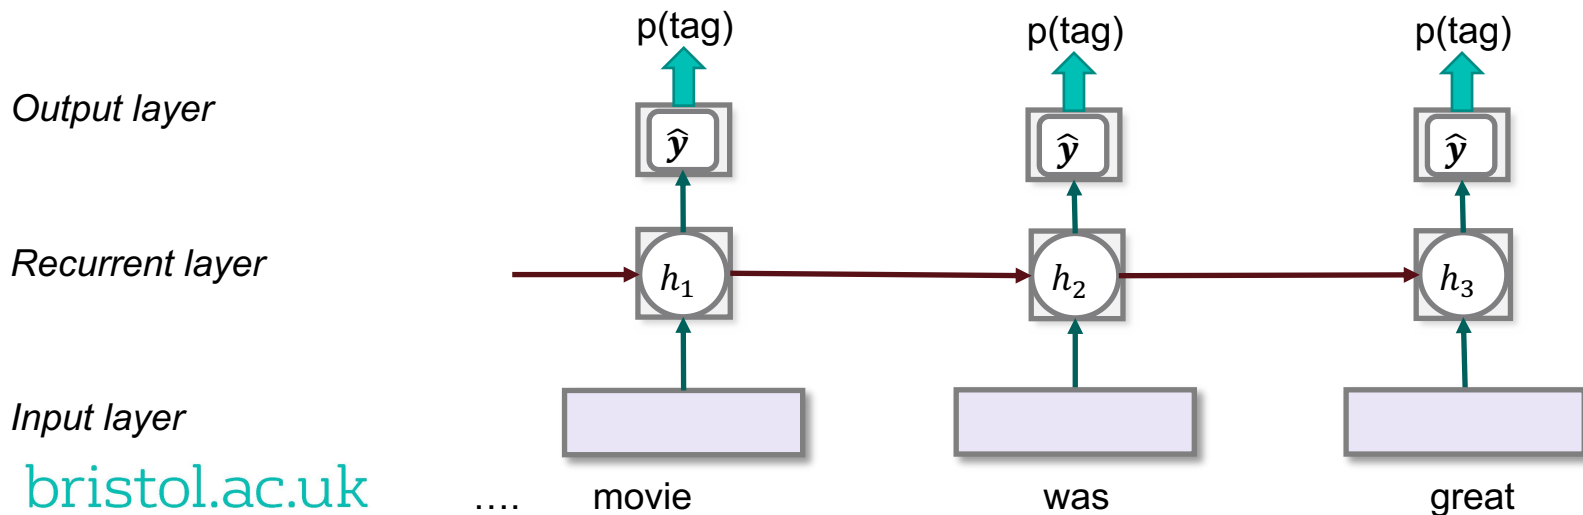
Fully-Connected (Dense) Layers

- Nodes in a fully-connected layer are connected to every node in the previous layer (here, this is the input sequence)
- The weights that are applied to each token depend solely on its position in the input sequence.



Recurrent Layers

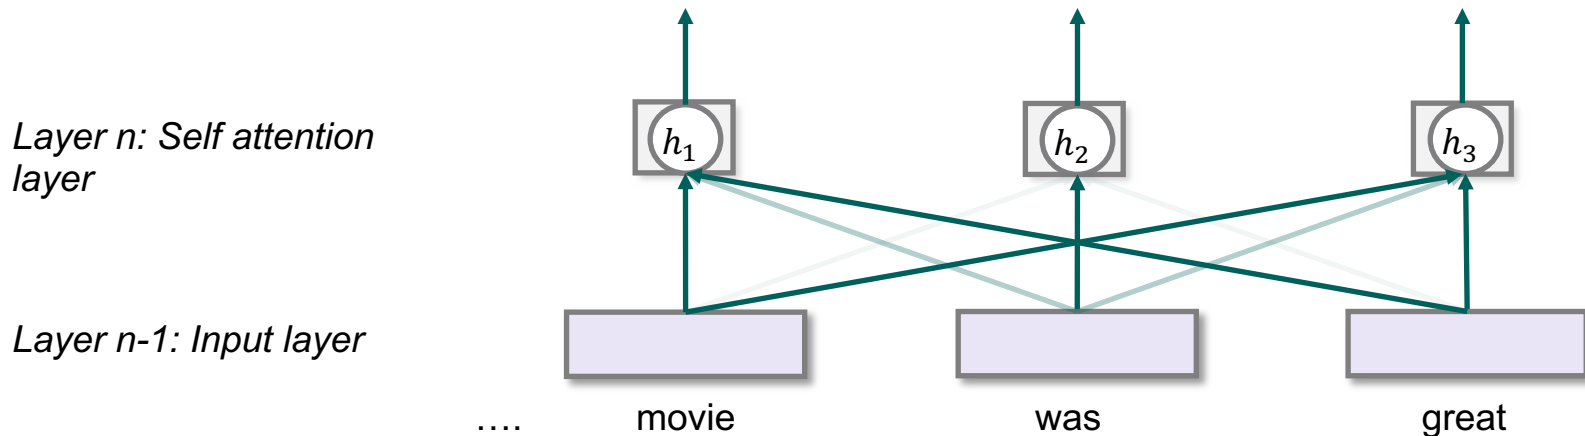
- Nodes in a recurrent layer are connected to the corresponding input and the output of the previous node in the sequence.
- All context information must pass along the sequence.



Connections To Anywhere in the Sequence?

- Ideally, the network would take **relevant context information from anywhere** in the sequence to process the current token.
- Connections would be created depending on the inputs.
- The weights applied to each input would then adapt to the inputs, rather than remaining fixed like a fully-connected layer.
- Direct connections mean shorter paths for context information to follow compared to RNNs.
- This can be achieved using **self-attention** layers.

Self-Attention



- **Attention** weights the connections between a sequence of values in layer $n-1$ and a sequence of nodes in layer n .
- The strengths of the connections are determined by comparing the values at each position in layer $n-1$ (the input sequence) with each other.


Self-Attention



How can we compare two word vectors?

- **Attention** weights the connections between a sequence of values in layer $n-1$ and a sequence of nodes in layer n .
- The strengths of the connections are determined by comparing the values at each position in layer $n-1$ (the input sequence) with each other.

Self-Attention

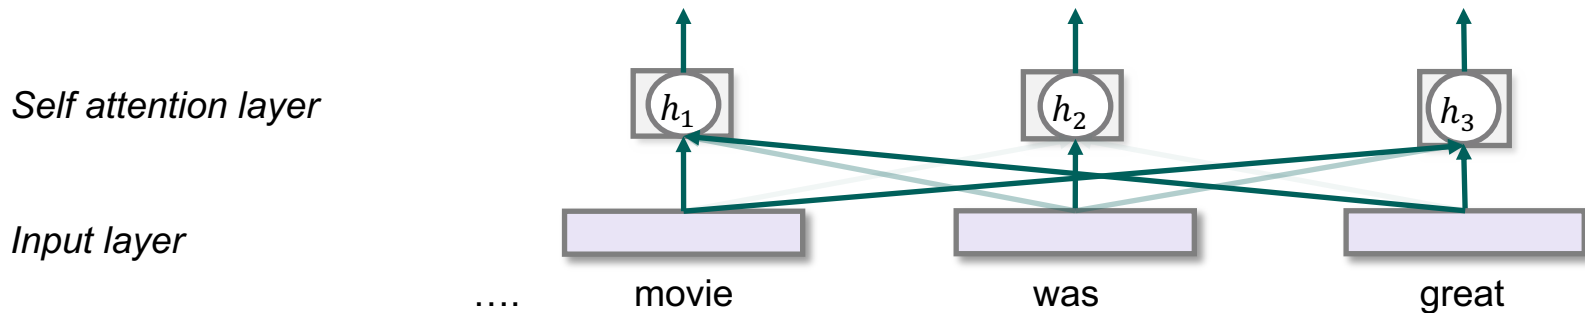


- Use dot product to compare the vectors for of two nodes in $n-1$:

$$score(x_i, x_j) = x_i \cdot x_j$$

- We also use dot product to compare vectors using cosine similarity.
- But this score is based only on the input vectors!
- It provides no way to learn how to make the connections between words in a sequence.

Self-Attention



- Use dot product to compare the vectors for of two nodes in $n-1$:

$$score(x_i, x_j) = W^q x_i \cdot W^k x_j$$

- W^q and W^k are matrices of weights that can be learned.
- Now we can learn how to words relate to each other in a sequence.


Self-Attention

- Scale and apply the softmax function to turn the score into attention weights that sum to 1:

$$\alpha(i, j) = \text{softmax}\left(\frac{\mathbf{W}^q \mathbf{x}_i \cdot \mathbf{W}^k \mathbf{x}_j}{\sqrt{d}}\right)$$

- $\alpha(i, j)$ indicates the proportion of attention we put on x_j in layer $n-1$ when determining the hidden state \mathbf{h}_i in layer n for token i .
- d is the number of dimensions of the embeddings in layer $n-1$

Self-Attention

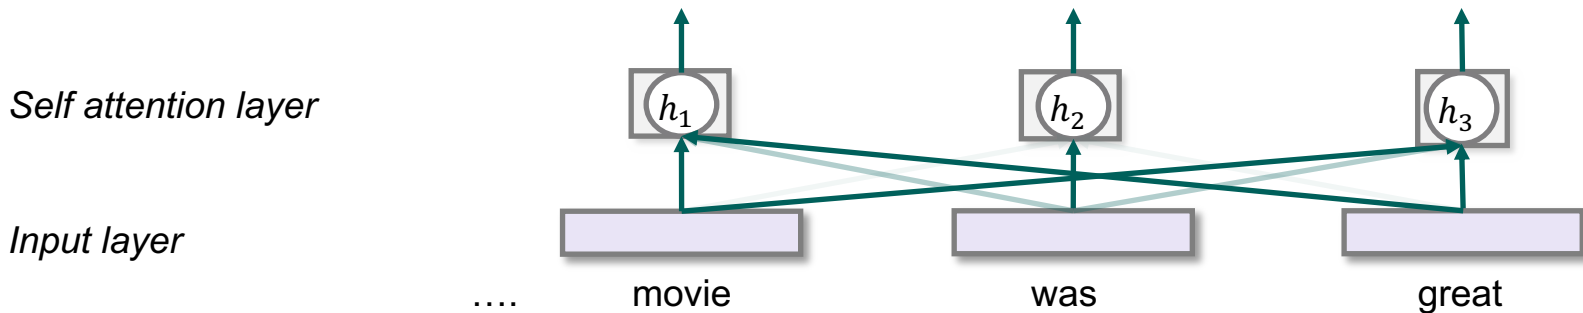


- Compute the hidden state \mathbf{h}_i for layer n as a weighted sum over the input sequence:

$$\mathbf{h}_i = \sum_{j=1}^N \alpha(i, j) \mathbf{W}^v \mathbf{x}_j$$

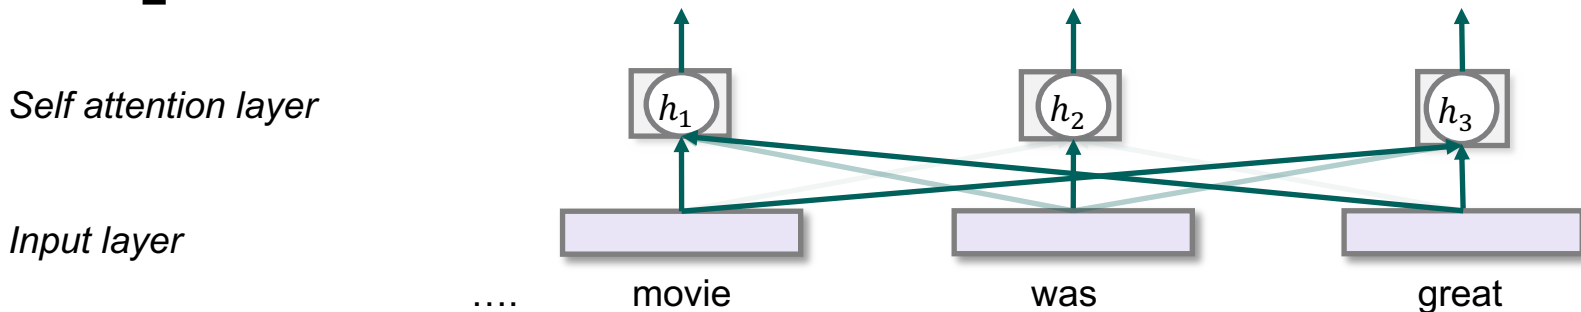
- \mathbf{W}^v is a matrix of learnable weights (equivalent to \mathbf{W} in the RNN)

Self-Attention



- Self-attention allows the model to select the relevant parts of the context when processing each token.
- Each hidden state h_i is computed independently, so can be computed in parallel.

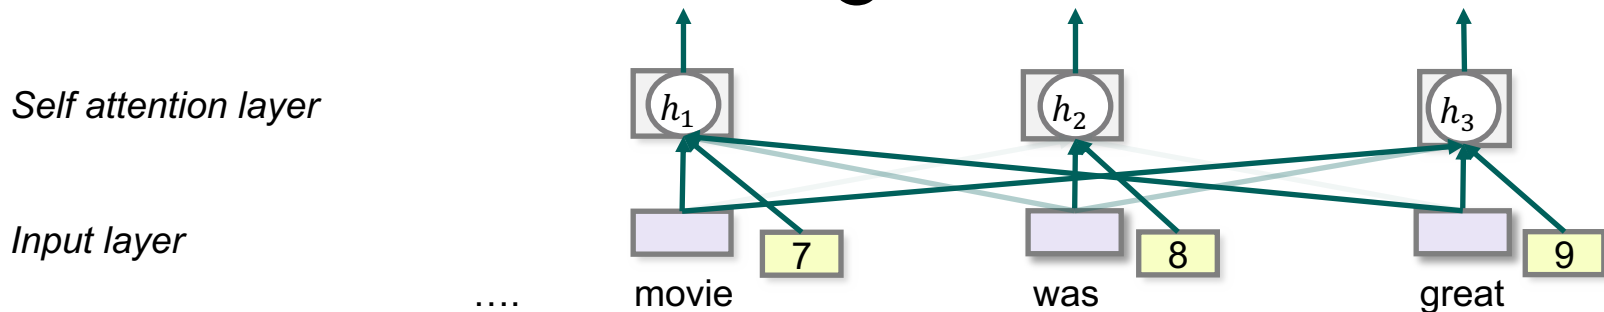
Sequential?



- The self-attention mechanism examines the whole sequence at once and doesn't know anything about word order...

If we ignore word order, what kind of model do we have?

Positional Embeddings



- The self-attention mechanism examines the whole sequence at once and doesn't know anything about word order...
- We're back to a bag of words!
- To provide information about word order, concatenate **positional embeddings** to the input embeddings.
- Positional embeddings are vectors that encode the position of each token in the sequence


Building a Model with Self-Attention

- In the NLP pipeline, we saw how successive processing steps extract syntactic and semantic information at different levels.
- A self-attention layer composes information at only one level of abstraction, so a single layer is not very powerful on its own.

How can we emulate this compositional behaviour with a neural network?

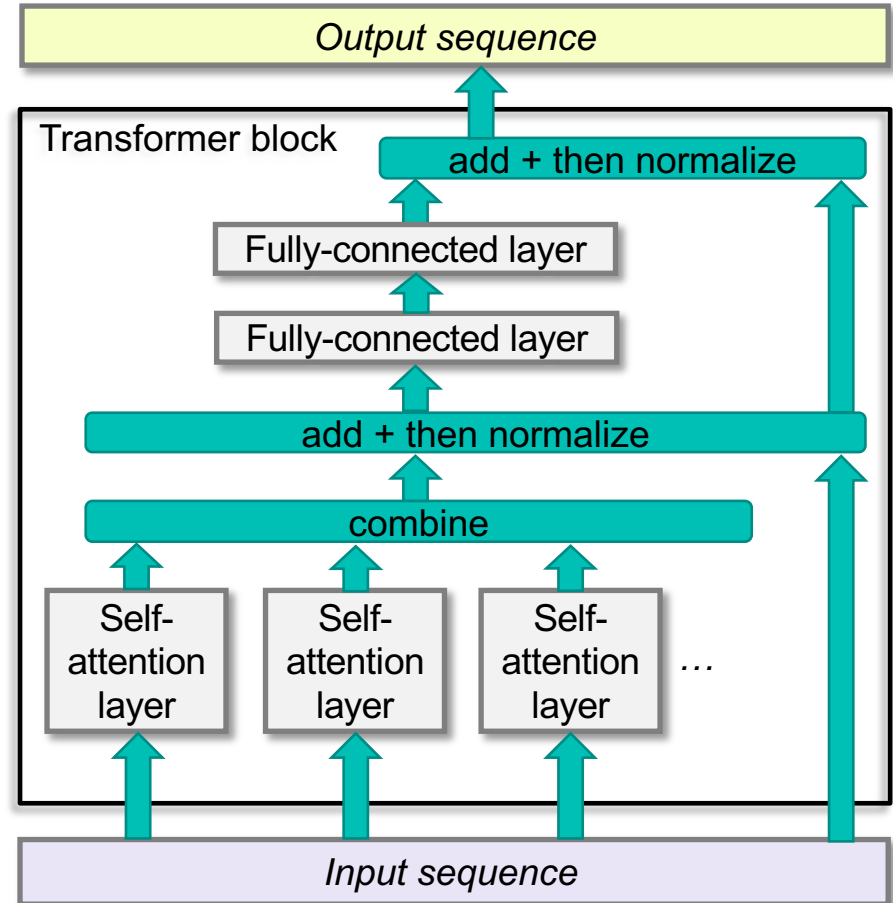
Transformers

- Stack multiple self attention layers to compose features at different levels of abstraction
- In a transformer, the self-attention layers are contained inside **transformer blocks**
- A transformer is a stack of transformer blocks →



Transformer Blocks

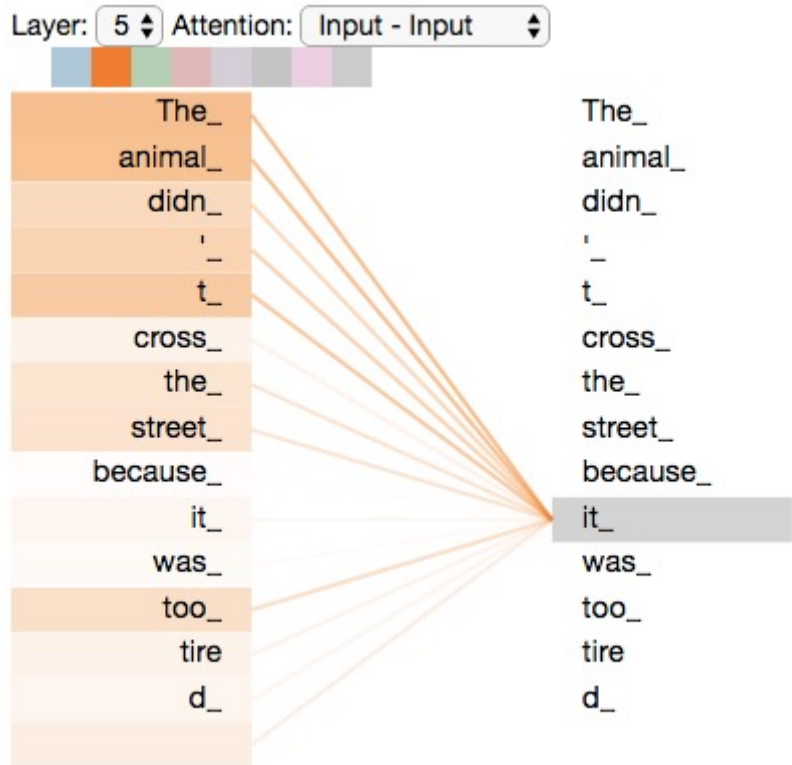
- Several self attention layers run in parallel: **multi-headed attention**
- Each attention head models a different kind of relation between the words in the input sequence.
- Fully-connected layers learn nonlinear functions



More on Transformers

- By avoiding recurrent connections, it is possible to train very large, deep transformers that excel at a wide range of language understanding and generation tasks
- More on transformers in the next video when we show how they can be used to provide contextualised embeddings
- See also:
 - [Attention is all you need, Vaswani et al., 2017](#)
 - [The illustrated Transformer](#)

Self-Attention Example

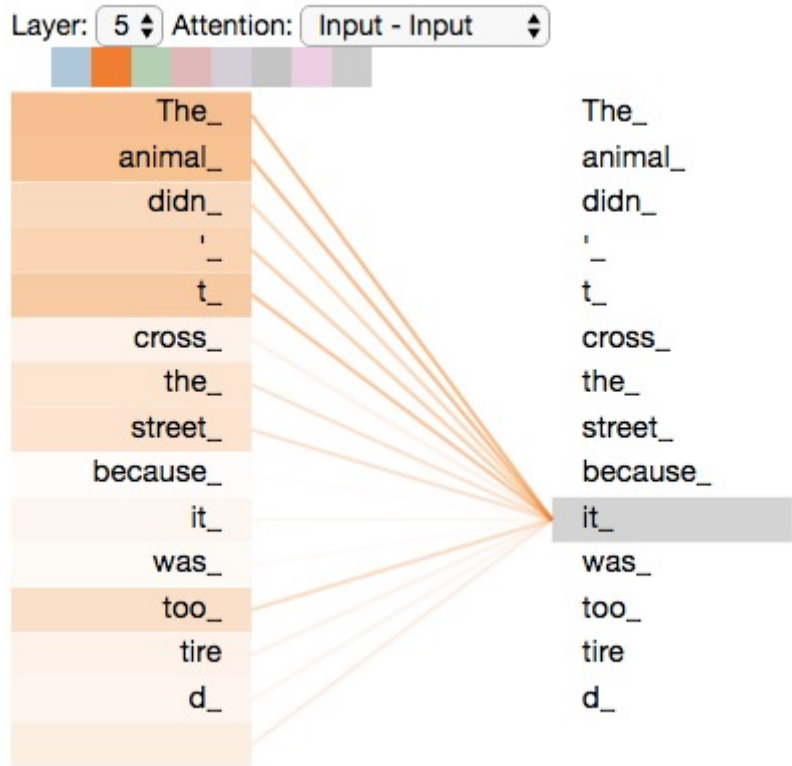


- The colours show the attention weights when processing the token 'it_'.

Can you think why certain input words have stronger attention weights for 'it_'?

Could this be useful in a task like relation extraction?

Self-Attention Example



- The colours show the attention weights when processing the token 'it_'.
- Strongest weights are on 'the animal'.
- Self-attention resolves 'it' to 'the animal' so the network can determine an encoding of 'it' that captures its meaning in this context.

Summary

- Parallel computation is not possible with RNNs and their handling of long-range dependencies is still limited.
- Self-attention layers allow the network to select the relevant context words from anywhere in the sequence and take a weighted sum of their values.
- Positional embeddings ensure that word order is considered.
- Transformers arrange many self-attention layers into large, deep models that extract different kinds of syntactic and semantic information and are highly effective in many tasks.